

Introduction to R

Computational Statistics

September 13, 2023

Generating long vectors

1a. The colon operator creates a sequence of integers via `start:end`. In R there is a more general sequence generator `seq()`. Use the help command `?seq` to learn about the function. Use `seq()` to create a sequence of numbers from 1 to 1001 in increments of 20.

Ans: Sequence of numbers from 1 to 1001 with an interval of 20 are as follows,

```
seq(from=1, to=1001, by=20)
```

```
## [1] 1 21 41 61 81 101 121 141 161 181 201 221 241 261 281
## [16] 301 321 341 361 381 401 421 441 461 481 501 521 541 561 581
## [31] 601 621 641 661 681 701 721 741 761 781 801 821 841 861 881
## [46] 901 921 941 961 981 1001
```

1b. The function `rep()` replicates the values in its input vector. Take a look at its help page. Explain the difference between `rep(1:4, times = 2)` and `rep(1:4, each = 2)`

Ans: `rep(1:4, times=2)` replicates the given sequence of 1 to 4 *as a whole* for two times as shown below,

```
rep(1:4, times=2)
```

```
## [1] 1 2 3 4 1 2 3 4
```

`rep(1:4, each=2)` replicates the **every element** of the sequence for 2 times i.e every integer in the sequence is repeated twice before printing the next integer.

```
rep(1:4, each=2)
```

```
## [1] 1 1 2 2 3 3 4 4
```

Functions and conditional statements

2a. Write a function `rescale01` that takes as input a numeric vector, and returns a rescaled vector whose components lie in $[0, 1]$. Linear scaling is applied to the finite numbers in the input vector, where the maximum and minimum are scaled to 1 and 0, respectively. Unexpected values should be handled as well. After the scaling, NA values should be unchanged, and positive infinity and negative infinity are set to 1 and 0, respectively. **Hint:** Use `is.infinite()`.

Ans: Let us consider a vector `v` with values as below,

```
v<-c(9,NA,5,3,Inf,-Inf,12:15)
v
```

```
## [1] 9 NA 5 3 Inf -Inf 12 13 14 15
```

Rescaling this vector while maintaining NA values will result in the following output,

```
v<-c(9,NA,5,3,Inf,-Inf,12:15)
rescale01<-function(l)
{
  RNG <- range(l,na.rm=TRUE, finite = TRUE) #excludes infinite values and ignores NA values
  lm<-(1-RNG[1])/(RNG[2]-RNG[1]) #rescaling using range function
  lm[lm == -Inf]<-0
  lm[lm == Inf]<-1
  lm
}
rescale01(v)
```

```
## [1] 0.5000000 NA 0.1666667 0.0000000 1.0000000 0.0000000 0.7500000
## [8] 0.8333333 0.9166667 1.0000000
```

The positive and negative infinity values can also be set to 1 and 0 respectively as mentioned above **or** can be changed after rescaling as below,

```
v<-c(9,NA,5,3,Inf,-Inf,12:15)
rescale01<-function(l)
{
  RNG <- range(l,na.rm=TRUE, finite = TRUE) #excludes infinite values and ignores NA values
  lm<-(1-RNG[1])/(RNG[2]-RNG[1])
  lm
}
p<-rescale01(v)
p[p == -Inf]<-0 #Replacing the negative infinity to 0
p[p == Inf]<-1 #Replacing the positive infinity to 1
p
```

```
## [1] 0.5000000      NA 0.1666667 0.0000000 1.0000000 0.0000000 0.7500000
## [8] 0.8333333 0.9166667 1.0000000
```

2b. With the `rescale01()` function, write a for loop to rescale every **numeric** column in the data frame `df`. **Hint:** Use if statement with `is.numeric()`.

```
set.seed(2023 - 9 - 13)
df <- data.frame(
  a = rnorm(5),
  b = 1:5,
  c = letters[6:10],
  d = c(-1, NA, Inf, -Inf, 1)
)
```

Ans: The given data frame will look as below,

```
set.seed(2023 - 9 - 13)
df <- data.frame(
  a = rnorm(5),
  b = 1:5,
  c = letters[6:10],
  d = c(-1, NA, Inf, -Inf, 1)
)
df
```

```
##           a b c      d
## 1  0.7019977 1 f    -1
## 2 -0.7778131 2 g     NA
## 3 -1.0240757 3 h    Inf
## 4  0.3027668 4 i   -Inf
## 5  0.3511394 5 j     1
```

Rescaling only the numeric columns in the above data frame can be achieved as below

```
set.seed(2023 - 9 - 13)
df <- data.frame(
  a = rnorm(5),
  b = 1:5,
  c = letters[6:10],
  d = c(-1, NA, Inf, -Inf, 1)
)
df
```

```
##           a b c      d
```

```
## 1 0.7019977 1 f -1
## 2 -0.7778131 2 g NA
## 3 -1.0240757 3 h Inf
## 4 0.3027668 4 i -Inf
## 5 0.3511394 5 j 1
```

```
rescale01<-function(l)
{
  RNG <- range(l,na.rm=TRUE, finite = TRUE) #excludes infinite values and ignores NA values
  lm<-(1-RNG[1])/(RNG[2]-RNG[1])
  lm[lm == -Inf]<-0
  lm[lm == Inf]<-1
  lm
}
outputdf<-data.frame(df) #preallocating the output space
for(i in seq_along(outputdf))
{
  if(is.numeric(outputdf[[i]])) #check if column is numeric
  {
    outputdf[[i]]<-rescale01(outputdf[[i]])
  }
}
outputdf
```

```
##          a    b c d
## 1 1.0000000 0.00 f 0
## 2 0.1426722 0.25 g NA
## 3 0.0000000 0.50 h 1
## 4 0.7687058 0.75 i 0
## 5 0.7967304 1.00 j 1
```

Note: If we observe, we set the values to positive and negative infinity to 1 and 0 respectively.

2c. Write a function `csum` that takes as input a finite positive integer n and computes the sum of cubes $S_n = 1^3 + 2^3 + \dots + n^3$. The function should include conditional checks. If an input is invalid (e.g., unexpected length or type, infinite, non-positive, not a number, etc.), the function should stop the execution with informative error message. Test your function with `csum(NA)`, `csum(NaN)`, `csum(Inf)`, `csum(1:3)`, `csum("two")`, `csum(0)`, `csum(10)`, `csum(-5L)`, `csum(1L)`, `csum(2L)`, `csum(3L)`.

Ans: A function that takes a finite positive integer and computes the sum of cubes is defined below,

```
csum <- function(n) {
  # Check if n is missing
  if (missing(n)) {
```

```

    stop("Input is missing (NA).")
  }
  #Check if the input is vector with length>1
  if ((length(n)>1) && (is.vector(n))) {
    stop("The input value is a vector and length is greater than 1")
  }
  # Check if n is NaN
  if (is.nan(n)) {
    stop("Input is not a number (NaN).")
  }
  # Check if n is infinite
  if (is.infinite(n)) {
    stop("Input is infinite.")
  }
  # Check if n is not a finite positive integer
  if (!is.numeric(n) || !is.finite(n) || n <= 0 || n != as.integer(n)) {
    stop("Input is not a finite positive integer.")
  }

  # Computing the sum of cubes
  Sn <- sum((1:n)^3)

  return(Sn)
}

```

When executed for multiple values of 'n' we get the below

```
csum(NA) #Testing the function for multiple values
```

```
## Error in csum(NA): Input is not a finite positive integer.
```

```
csum(NaN)
```

```
## Error in csum(NaN): Input is not a number (NaN).
```

```
csum(Inf)
```

```
## Error in csum(Inf): Input is infinite.
```

```
csum(1:2)
```

```
## Error in csum(1:2): The input value is a vector and length is greater than 1
```

```
csum("two")
```

```
## Error in csum("two"): Input is not a finite positive integer.
```

```
csum(0)
```

```
## Error in csum(0): Input is not a finite positive integer.
```

```
csum(10)
```

```
## [1] 3025
```

```
csum(-5L)
```

```
## Error in csum(-5L): Input is not a finite positive integer.
```

```
csum(1L)
```

```
## [1] 1
```

```
csum(2L)
```

```
## [1] 9
```

```
csum(3L)
```

```
## [1] 36
```

2d. Write a while loop to print the value of S_n for $n = 1, 2, \dots$ until $S_n > 2000$. At each iteration, you can print the sum of cubes using an R command similar to this: `print(paste0("n = ", n, ", sum of cubes is ", csum(n)))`.

Ans: Let us initialize the sum of cubes to be zero. Now, printing the sum of cubes until the sum reaches 2000 is defined via the below code and will result in the following output.

```
csum<-function(num){  
  sum((1:num)^3)  
}  
num<-1  
while(csum(num)<=2000){  
  print(paste0("n= ", num,", sum of cubes is ", csum(num)))  
  num<-num+1  
}
```

```
## [1] "n= 1, sum of cubes is 1"
## [1] "n= 2, sum of cubes is 9"
## [1] "n= 3, sum of cubes is 36"
## [1] "n= 4, sum of cubes is 100"
## [1] "n= 5, sum of cubes is 225"
## [1] "n= 6, sum of cubes is 441"
## [1] "n= 7, sum of cubes is 784"
## [1] "n= 8, sum of cubes is 1296"
```

Conversion of abbreviations into full words

A character vector `fall_abbr` stores information of fall months as `sept`, `oct` and `nov`:

```
fall_abbr <- sample(c("sept", "oct", "nov"), size = 30, replace = TRUE)
fall_abbr
```

```
## [1] "sept" "sept" "sept" "nov" "sept" "nov" "nov" "nov" "sept" "sept"
## [11] "oct" "oct" "sept" "nov" "sept" "sept" "oct" "nov" "sept" "oct"
## [21] "nov" "oct" "oct" "nov" "nov" "oct" "oct" "nov" "sept" "sept"
```

We will use different ways to convert the abbreviations into full words: `"sept" → "September"`, `"oct" → "October"`, and `"nov" → "November"`.

3a. Use a for loop and the switch statement to convert all the elements of `fall_abbr` into full words, and store the results in a vector named `fall_full`.

Ans: After the conversion, the `fall_full` will result in,

```
fall_full<-character(length(fall_abbr))
for(i in seq_along(fall_abbr)){
  abbr<-fall_abbr[i]
  full_name<-switch(abbr,
    sept="September",
    oct="October",
    nov="November",
    stop("Invalid value")
  )
  fall_full[i]<-full_name
}
fall_full
```

```
## [1] "September" "September" "September" "November" "September" "November"
## [7] "November" "November" "September" "September" "October" "October"
## [13] "September" "November" "September" "September" "October" "November"
## [19] "September" "October" "November" "October" "October" "November"
## [25] "November" "October" "October" "November" "September" "September"
```

3b. Carry out the conversion by creating a lookup table and subsetting the lookup table with characters.

Ans: Creating a lookup table and referencing this data to the fall_abbr table will give us the desired result.

```
lookup_table<-c(sept="September", oct="October", nov="November")
fall_full<-lookup_table[fall_abbr]
fall_full
```

```
##      sept      sept      sept      nov      sept      nov
## "September" "September" "September" "November" "September" "November"
##      nov      nov      sept      sept      oct      oct
## "November" "November" "September" "September" "October" "October"
##      sept      nov      sept      sept      oct      nov
## "September" "November" "September" "September" "October" "November"
##      sept      oct      nov      oct      oct      nov
## "September" "October" "November" "October" "October" "November"
##      nov      oct      oct      nov      sept      sept
## "November" "October" "October" "November" "September" "September"
```

3c. The function ifelse() is a vectorized version of the if-else statements. Use the function to convert all the elements of fall_abbr.

Ans: **Vectorized version**

```
fall_full <- ifelse(fall_abbr %in% names(lookup_table), lookup_table[fall_abbr], "Unknown Fall M
fall_full
```

```
## [1] "September" "September" "September" "November" "September" "November"
## [7] "November" "November" "September" "September" "October" "October"
## [13] "September" "November" "September" "September" "October" "November"
## [19] "September" "October" "November" "October" "October" "November"
## [25] "November" "October" "October" "November" "September" "September"
```